

# Awkward + CCCL

Presentation by: Maksym Naumchyk

Supervisor: Ianna Osborne

Thanks to  NVIDIA® team for help and support!

# How does it work?

# CUDA kernels

```
template <typename T, typename C, typename U>
__global__ void
awkward_reduce_sum_a(
    T* toptr,
    const C* fromptr,
    const U* parents,
    int64_t lenparents,
    int64_t outlength,
    T* temp,
    uint64_t invocation_index,
    uint64_t* err_code) {
    if (err_code[0] == NO_ERROR) {
        int64_t thread_id = blockIdx.x * blockDim.x + threadIdx.x;

        if (thread_id < outlength) {
            toptr[thread_id] = 0;
        }
    }
}

template <typename T, typename C, typename U>
__global__ void
awkward_reduce_sum_b(
    T* toptr,
    const C* fromptr,
    const U* parents,
    int64_t lenparents,
    int64_t outlength,
    T* temp,
    uint64_t invocation_index,
    uint64_t* err_code) {
    if (err_code[0] == NO_ERROR) {
        int64_t idx = threadIdx.x;
        int64_t thread_id = blockIdx.x * blockDim.x + idx;

        if (thread_id < lenparents) {
            temp[thread_id] = fromptr[thread_id];
        }
        __syncthreads();
    }
}
```

and more ...

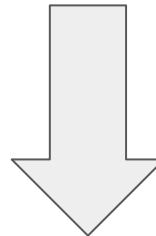
The whole sum() operation using CUDA kernels:

[https://github.com/scikit-hep/awkward/blob/main/src/awkward/\\_connect/cuda/cuda\\_kernels/awkward\\_reduce\\_sum.cu](https://github.com/scikit-hep/awkward/blob/main/src/awkward/_connect/cuda/cuda_kernels/awkward_reduce_sum.cu)

# CCCL

We can feed ListOffsetArrays directly to the CCCL algorithms:

```
<ListOffsetArray len='6'>
    <offsets><Index dtype='int64' len='7'>
        [0 1 3 5 8 8 9]
    </Index></offsets>
    <content><NumpyArray dtype='int64' len='9'>[1 2 3 4 5 6 7 8 9]</NumpyArray></content>
</ListOffsetArray>
```



```
segmented_reduce(
    input_data, output, start_offsets, end_offsets, sum_operation, h_initial, n_segments
)
```

Why do we need this?

# Performance improvement!

ak.sum()

— CPU    — CUDA kernels    — CCCL



ak.sum()

File size

128 B

280.0 MB

2520.0 MB

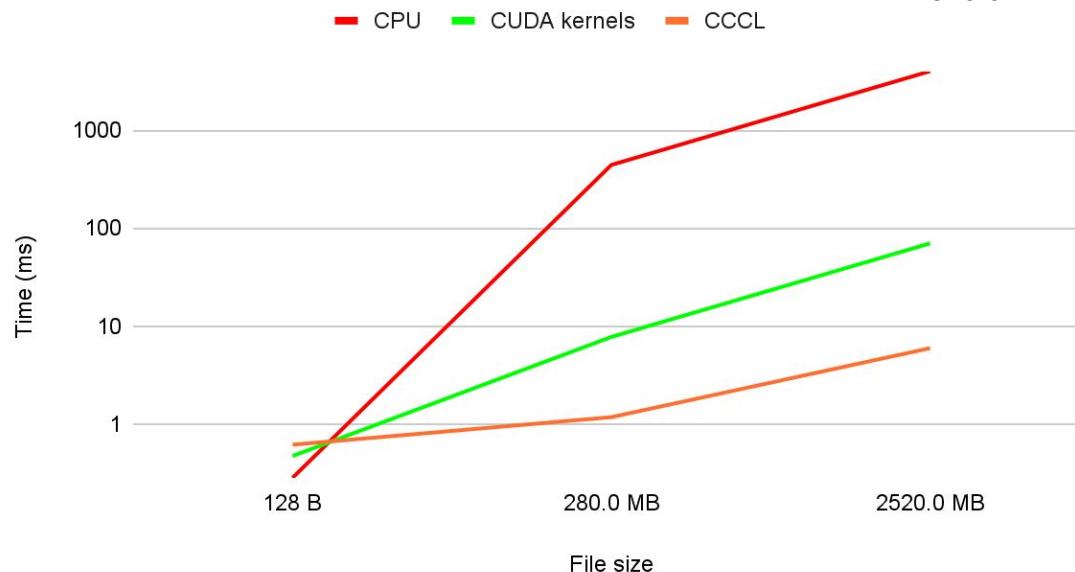
CPU

CUDA kernels

CCCL

# Performance improvement!

ak.argmax()



ak.argmax()

CPU

CUDA kernels

CCCL

(using  
unary\_transform)

File size

128 B

280.0 MB

2520.0 MB

0.284 ms

442 ms

4000 ms

0.474 ms

7.77 ms

70.0 ms

0.618 ms

1.18 ms

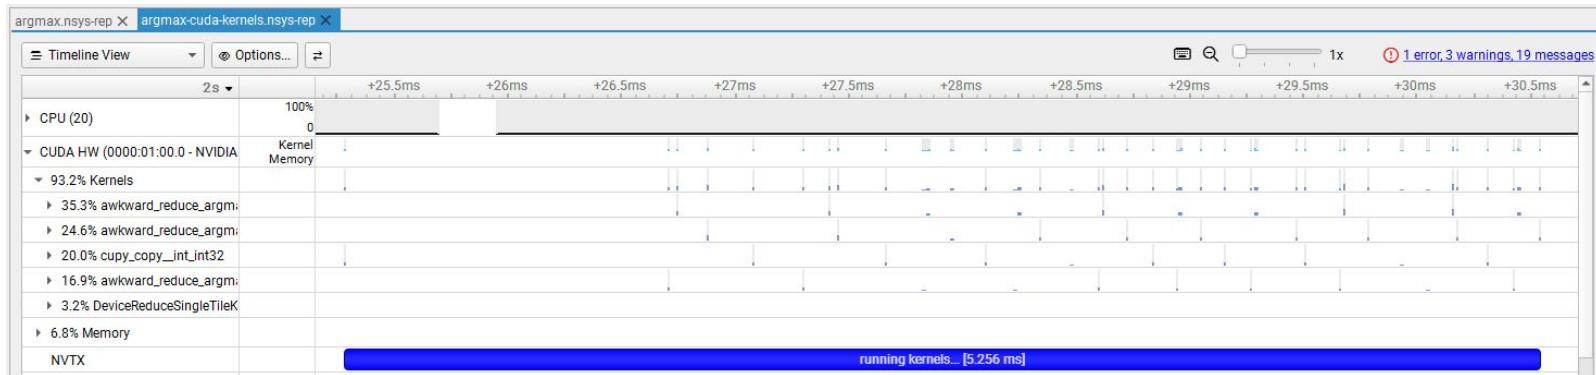
5.97 ms

Even more improvements!

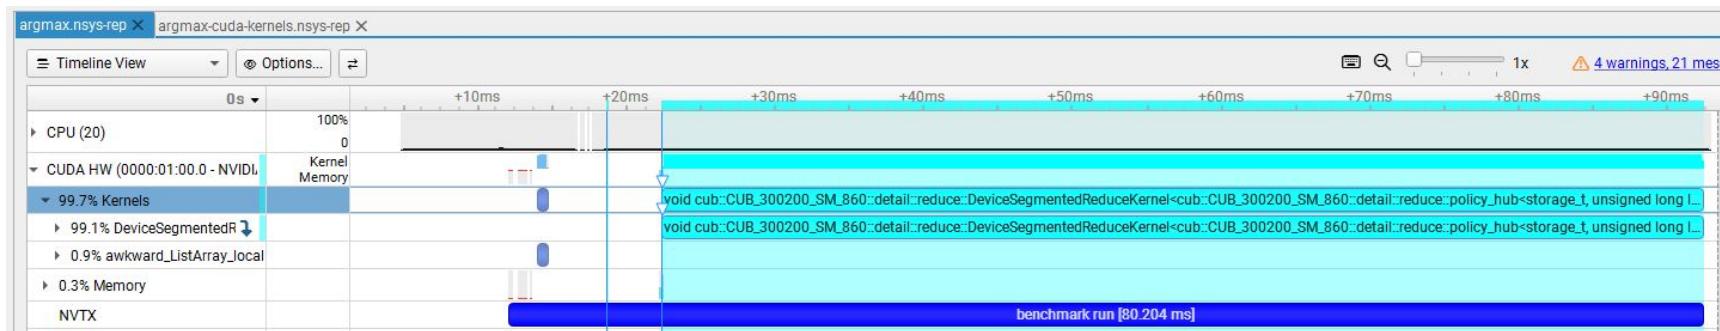
ak.sort() with cccl

ak.combinations() using  
cp.searchsorted

# Kernels fusion



A lot of small CUDA kernels (previously with awkward) vs →



one big kernel fused with CCCL

# Going a step further

# Introducing (new) kernels

```
# select only electrons with pt > 40
selected_electrons = events.electrons[events.electrons.pt > 40.0]

def cond_muon(x):
    return (x[0] > 20.0) & (abs(x[1]) < 2.4)
selected_muons = filter_lists(events.muons, cond_muon)

def filter_lists(array, cond):
    it, meta = awkward_to_cccl_iterator(array)
    in_segments = meta["offsets"]
    out_array = empty_like(array)
    it_out, meta_out = awkward_to_cccl_iterator(out_array)
    out_segments = meta_out["offsets"]
    num_items = meta["count"]
    segmented_select(
        it,
        in_segments,
        it_out,
        out_segments,
        cond,
        num_items
    )
    return reconstruct_with_offsets(out_array, out_segments)
```

LHC data analysis playground by Peter: <https://github.com/scikit-hep/awkward/pull/3734/files>

Playground using CCCL by Ashwin: <https://github.com/scikit-hep/awkward/blob/main/studies/cccl/playground.py>

# [Prototype] Intermediate Representation (IR)

```
selected_muons = filter_lists(events.muons, cond_muon)
two_muons = select_lists(
    selected_muons, (list_sizes(selected_muons) == 2).astype('int8'))
```

---

```
class IsEqual(IRNode):
    def lower(self, a, b):
        return a == b

selected_muons = Filter(events.muons, cond_muon)
two_muons = SelectLists(
    selected_muons, (IsEqual(ListSizes(selected_muons), 2)))
```

# [Prototype] IR wrapper

```
arr = ak.Array([[1, 2, 3], [4, 5], [6, 7, 8, 9]], backend="cuda")

# Create Lazy wrapper
lazy_arr = ak.cuda.lazy(arr)

# Build computation graph without executing
transformed = lazy_arr * 2 + 1
result = transformed.filter(lazy_arr > 3)

# Execute the computation
print("Computed result:")
print(result.compute())
```

```
Computed result:
[], [9, 11], [13, 15, 17, 19]]
```

Github issue on this prototype: <https://github.com/scikit-hep/awkward/pull/3792>  
Feel free to discuss it there!

# Thank *you* for attention!

Useful links:

NVIDIA CCCL: <https://nvidia.github.io/cccl/python/>

Ashwin's presentation on CCCL for HEP: <https://indico.cern.ch/event/1566263/contributions/6733091/>  
[[youtube](#) recording]